



WEAVE - SENSIBLE YARN

Nitin Motgi & Terence Yim

Continuity Proprietary and Confidential

INTRODUCTION





- **Who we are ?**

- Group of techies who have experienced Big Data in a variety of ways.
- We understand the pain of building Big Data Applications - we have lived through it.

- **What do we do ?**

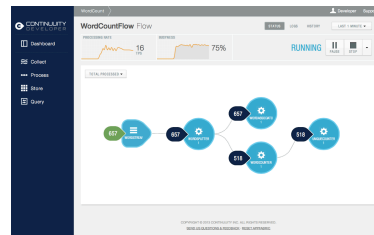
- **Mission** : *Democratize the development of the next generation of Big Data applications.*
- Building the industry's first Big Data Application Fabric(TM)

- **Products**

-  Developer Suite (public beta available now)
 - Downloadable developer edition. Includes SDK, tutorial, samples, docs, etc
-  Developer Sandbox
 - Free, Self-Service, hosted cloud environment
-  Virtual Private Cloud
 - Single tenant, private cloud PaaS. Fully Integrated BigData Application Platform
-  On-Premise



BIG FLOW



What is BigFlow?

- A core Processor; a real-time stream processing engine (Similar to Storm and S4)
- BigFlow provides
 - Tight integration with your existing Hadoop/HBase cluster
 - Exactly-once execution vs. At-least-once execution (your apps do not have to be idempotent)
 - Supports transactions for consistency
 - Direct integration with a data storage engine (Continuity Data Fabric) and Datasets
 - Beautiful User interface for application management
 - Utilizes YARN for deployment and supports runtime elastic scalability of Flows



DEMO

Continuity Proprietary and Confidential





YARN PAIN POINTS

- YARN is not for novice Big Data programmers
 - High ramp-up time
- Lot of boiler plate code to build simple application
 - 80% of applications are generally simple
- YARN is built around applications that terminate
 - Logs are available only on completion of application
- No standard support for
 - Application lifecycle management
 - Communication between Container(s) & Application Master
 - Handling Application level errors





MOTIVATION

- Running in YARN should be as simple as running Threads in Java
 - Not designed to be a replacement for YARN
- Designed to simplify building, debugging & running of YARN Applications
 - A simple programming model
- Hide messy details of YARN
 - Simplified API for specifying, running & managing an application
 - Simplified way to specify and manage stage(s) of a Application lifecycle
 - Generic Application Master to better support simple applications
 - Simpler archive management
 - Better control over application errors



SIMPLICITY

Distributed Shell - Vanilla YARN

~ 500 lines

```
140. public class ApplicationMaster {
141.
142.     private ANHMCClientAsync resourceManager;
143.
144.     public boolean run() throws YarnRemoteException {
145.         LOG.info("Starting ApplicationMaster");
146.
147.     }
148.
149.     public void finish() {
150.         for(Thread launchThread: launchThreads) {
151.             ...
152.         }
153.         FinalApplicationStatus appStatus;
154.         String appMessage = null;
155.     }
156.
157.     public ContainerRequest setupContainerAsksForRM(int numContainers) {
158.         Resource capability = Resource.newInstance(1, 1);
159.         capability.setMemory(containerMemory);
160.
161.         ContainerRequest request = new ContainerRequest(capability, ...);
162.     }
163.
164.     public class Client extends YarnClientImpl {
165.
166.     public boolean run() throws IOException {
167.         GetNewApplicationResponse newApp = super.getNewApplication();
168.         ApplicationId appId = newApp.getApplicationId();
169.
170.         return monitorApplication(appId);
171.     }
172.
173.     private boolean monitorApplication(ApplicationId appId) throws ... {
174.         while(true) {
175.             ApplicationReport report = super.getApplicationReport(appId);
176.
177.         }
178.
179.         private void forceKillApplication(ApplicationId appId) throws ... {
180.             super.killApplication(appId);
181.
182.         }
183.     }
184. }
```

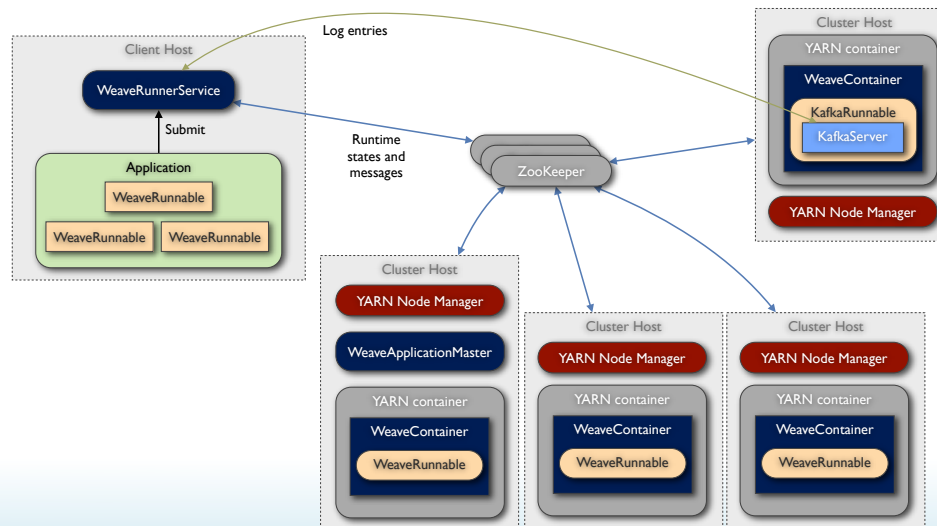
Distributed Shell - WEAVE

~ 50 lines

```
201. public class DistributedShell extends AbstractWeaveRunnable {
202.     static Logger LOG = LoggerFactory.getLogger(DistributedShell.class);
203.     private String command;
204.
205.     public DistributedShell(String command) {
206.         this.command = command;
207.     }
208.
209.     @Override
210.     public WeaveRunnableSpecification configure() {
211.         return WeaveRunnableSpecification.builder().with()
212.             .setName("Executor " + command)
213.             .withArguments(ImmutableMap.of("cmd", this.command))
214.             .build();
215.     }
216.
217.     @Override
218.     public void initialize(WeaveContext context) {
219.         this.command = context.getSpecification().getArguments("cmd");
220.     }
221.
222.     @Override
223.     public void run() {
224.         try {
225.             Process process = Runtime.getRuntime().exec(this.command);
226.             InputStream stdin = proc.getInputStream();
227.             BufferedReader br = new BufferedReader(new InputStreamReader(stdin));
228.
229.             String line;
230.             while( (line = br.readLine()) != null) {
231.                 LOG.info(line);
232.             }
233.         } catch (Throwable t) {
234.             LOG.error(t.getMessage());
235.         }
236.     }
237.
238.     public static main(String[] args) {
239.         String zkStr = args[0];
240.         String command = args[1];
241.
242.         yarnConfig = ...
243.         WeaveRunnerService weaveService = new YarnWeaveRunnerService(yarnConfig, zkStr);
244.         weaveService.startShell();
245.         WeaveController controller = weaveService.prepare(new DistributedShell(command))
246.             .addLogHandler(new PrinterLogHandler(new PrintWriter(System.out)))
247.             .start();
248.     }
```



ARCHITECTURE



RUNNABLE

```
// Defines a Task
public class DistributedShell implements WeaveRunnable
{
    static Logger LOG = ...
    private String cmd;

    //....

    @Override
    public void run() {
        try {
            new ProcessBuilder(cmd).start().waitFor();
        } catch (Exception e) {
            LOG.error(e.getMessage(), e);
        }
    }
}
```

- ▶ Class implements WeaveRunnable is a single task.
- ▶ SLF4J logger can be used for logging
 - ▶ logs are collected and sent back to client using Kafka
 - ▶ An extra container is started to run Kafka server.
- ▶ WeaveRunnable can be run in
 - ▶ Thread
 - ▶ YARN Container
- ▶ Easy! - Piece of cake.



RUNNER

```
// Starts a task
WeaveRunner runner = ...;
WeaveController controller =
    runner.prepare(new DistributedShell("ls -al"))
        .addLogHandler(new PrinterLogHandler(
            new PrintWriter(System.out)
        ))
        .start();
//...
controller.sendCommand(Commands.create("flush"));
//...
controller.stop();
```

- WeaveRunner
 - Can run Runnable or Application
 - Packages the class and it's dependencies along with additional resources specified
 - Attaches log collector for collecting all the logs from containers
- WeaveController is used to manage the lifecycle of a runnable or application
- Controller also provides ability to send commands to running containers



APPLICATION

```
// Defines an Application
public class WebServer implements WeaveApplication {
    @Override
    public WeaveSpecification configure() {
        return WeaveSpecification.Builder.with()
            .setName("Jetty WebServer")
            .withRunnables()
                .add(new JettyWebServer())
            .withLocalFiles()
                .add("html-pages.tgz", pages, true)
            .apply()
            .add(new LogsCollector())
        .anyOrder()
        .build();
    }
}
```

- WeaveApplication defines a collection of runnables and their behavior
- Application is specified by a WeaveSpecification
- Application can specify any additional local directories to be made available for container to run.
- Weave starts containers in no order.



APPLICATION

```
// Defines Application with specific tasks ordering
public class DataApplication implements WeaveApplication {
    @Override
    public WeaveSpecification configure() {
        return WeaveSpecification.Builder.with()
            .setName("Cool Data Application")
            .withRunnables()
                .add("reader", new InputReader())
                .add("processor", new Processor())
                .add("writer", new OutputWriter())
            .order()
                .first("reader")
                .next("processor")
                .next("writer")
            .build();
    }
}
```



DEMO

Continuity Proprietary and Confidential



ROADMAP

- First release at the end of April
- Features for upcoming releases (in no order)
 - API support for managing lifecycle hooks of application & it's containers
 - Extendable Application Master
 - Support for command dispatching
 - Ability to programmatically control resource allocation
 - Advanced error handling capabilities
 - E.g. Fail application if 50% of my containers fail
 - E.g. Fail application if it take more 'x' seconds to allocated containers for application
 - Contribs
 - Running HBase using Weave
 - Running Jetty HTTP Server using Weave
 - Support for application level metrics
 - WeaveRunnerServer for managing multiple application
 - Better UI
 - Ability to remote debug a running container
 - LXC & QoS
 - Improve! APIs, Improve! APIs
 - . . .



THANK YOU

git@github.com:continuuity/weave.git

